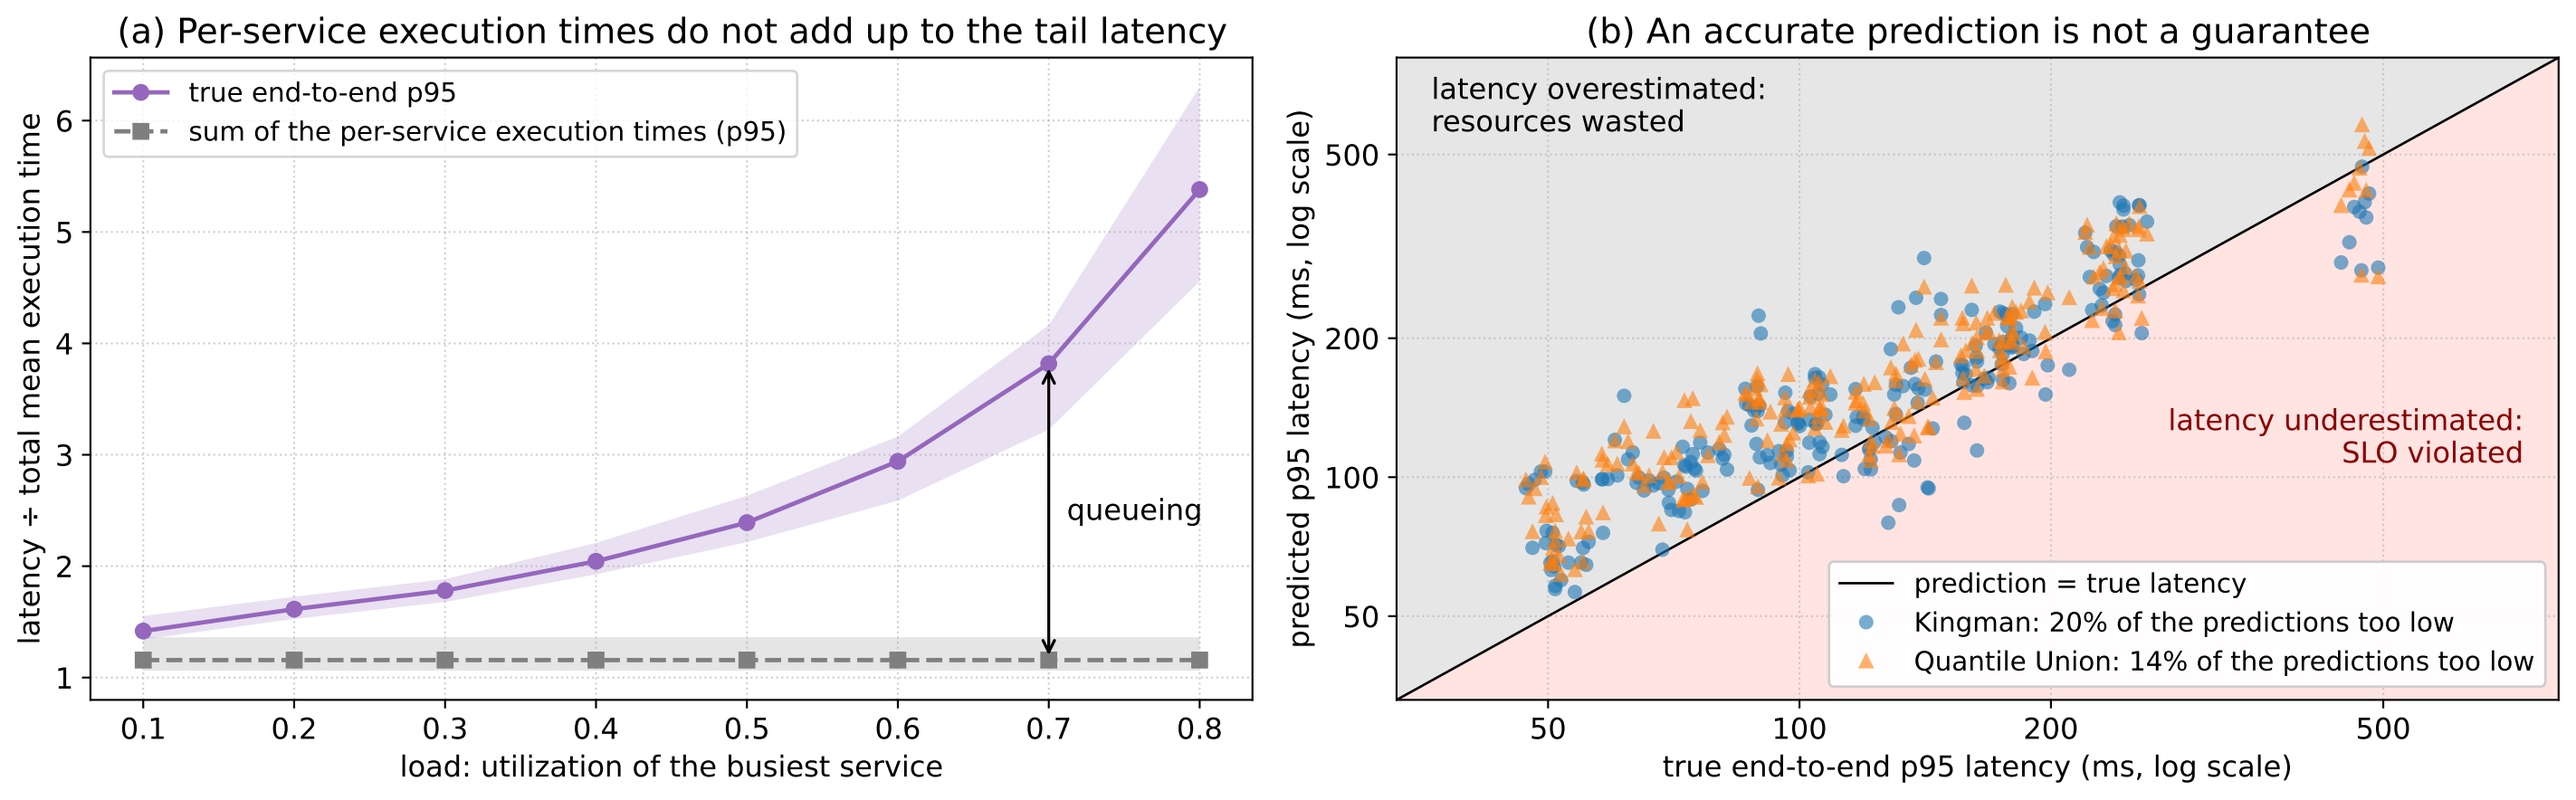
(a) Per-service execution times do not add up to the tail latency
true end-to-end p95
sum of the per-service execution times (p95)
queueing
latency ÷ total mean execution time
load: utilization of the busiest service
(b) An accurate prediction is not a guarantee
latency overestimated:
resources wasted
latency underestimated:
SLO violated
prediction = true latency
Kingman: 20% of the predictions too low
Quantile Union: 14% of the predictions too low
predicted p95 latency (ms, log scale)
true end-to-end p95 latency (ms, log scale)