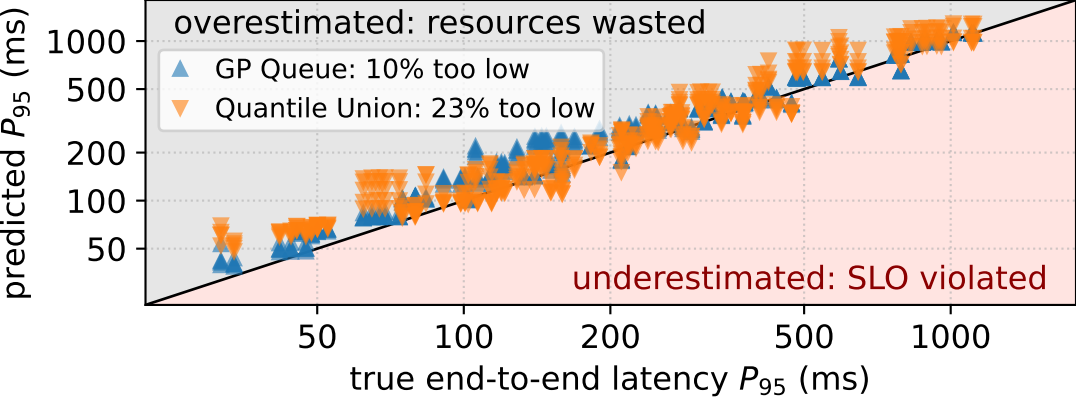
overestimated: resources wasted
GP Queue: 10% too low
Quantile Union: 23% too low
underestimated: SLO violated
predicted P_{95} (ms)
1000
500
200
100
50
50
100
200
500
1000
true end-to-end latency P_{95} (ms)