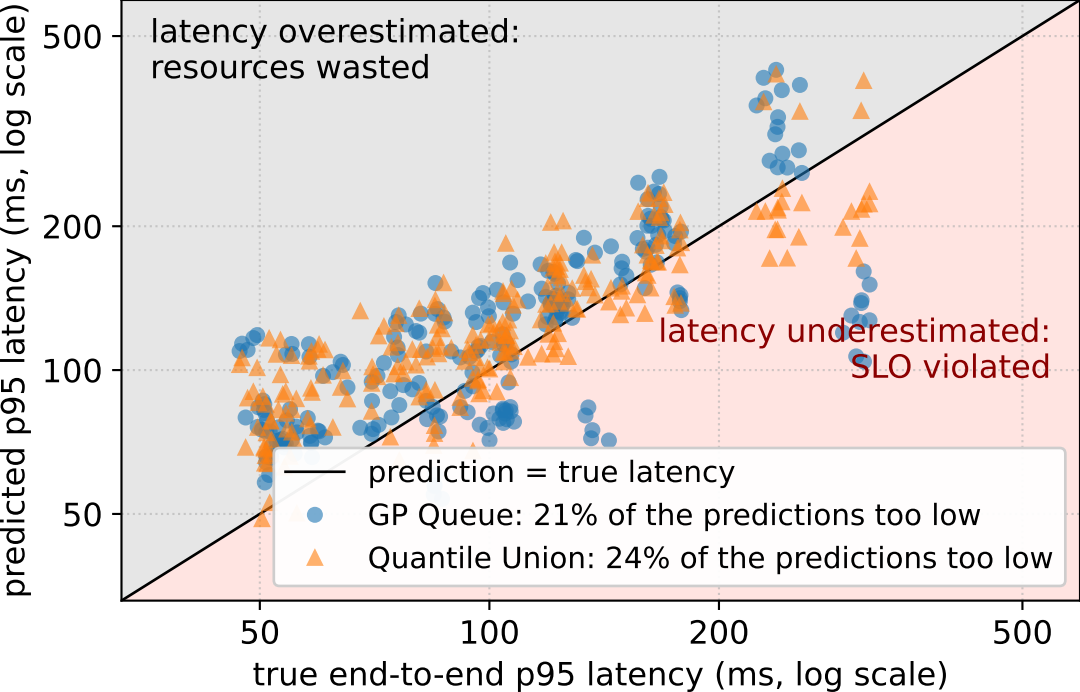
latency overestimated:
resources wasted
latency underestimated:
SLO violated
prediction = true latency
GP Queue: 21% of the predictions too low
Quantile Union: 24% of the predictions too low
predicted p95 latency (ms, log scale)
true end-to-end p95 latency (ms, log scale)
50
100
200
500
50
100
200
500