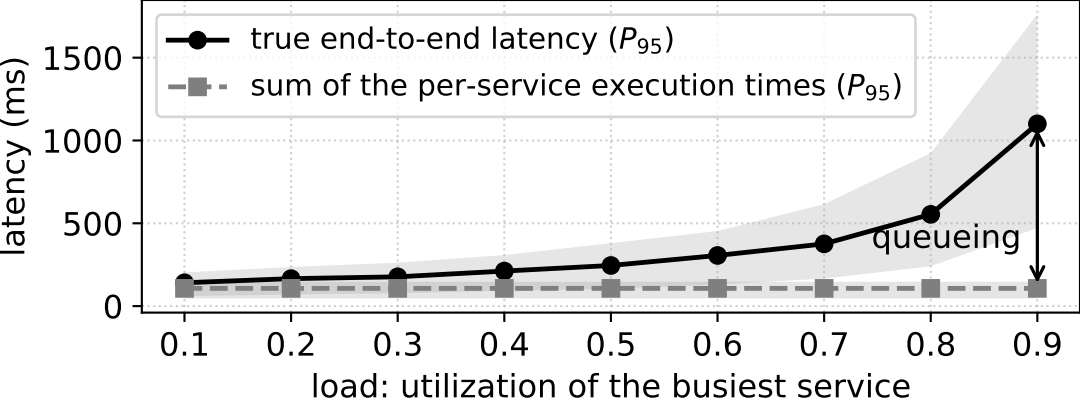
true end-to-end latency (P_{95})
sum of the per-service execution times (P_{95})
queueing
latency (ms)
0
500
1000
1500
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.8
0.9
load: utilization of the busiest service